



UPPSALA UNIVERSITET

Monte Carlo Computations

Yuanjing Yang

June 27, 2025

1 Introduction

This project implements a MPI parallel version of SSA algorithm to simulate the spread of malaria epidemic, which is primarily through the bite of infected female Anopheles mosquitoes. With parallel computing, Monte Carlo simulations can be independent and executed concurrently, enabling efficient execution.

2 Algorithm

2.1 Sequential Algorithm

The sequential component is based on the SSA algorithm, which simulating the time evolution of a system with probabilistic reaction events. In this simulation:

- (1) The state of the system is represented by a 7-dimensional integer vector, describing different compartments of human and mosquito populations.
- (2) The state evolves over time, resulting 15 reactions.
- (3) At each step, the algorithm calculates reaction propensities, determines the time to the next event, and selects which reaction occurs.

SSA Algorithm [1]:

Algorithm 1 Gillespie's direct method (SSA)

- 1: Set a final simulation time T , current time $t = 0$, initial state $\mathbf{x} = \mathbf{x}_0$
 - 2: **while** $t < T$ **do**
 - 3: Compute $\mathbf{w} = \text{prob}(\mathbf{x})$
 - 4: Compute $a_0 = \sum_{i=1}^R w(i)$
 - 5: Generate two uniform random numbers $u_1, u_2 \in (0, 1)$
 - 6: Set $\tau = -\ln(u_1)/a_0$
 - 7: Find r such that $\sum_{k=1}^{r-1} w(k) < a_0 u_2 \leq \sum_{k=1}^r w(k)$
 - 8: Update the state vector $\mathbf{x} = \mathbf{x} + P(r, :)$
 - 9: Update time $t = t + \tau$
 - 10: **end while**
-

2.2 Parallel Algorithm

Here I used master-worker patterns [3] with dynamic load balancing [2], the procedures is below:

1. Master process(Rank 0) distributes simulation tasks to workers.
2. worker process receive task assignments from master and execute SSA simulations independently.

MPI Function used:

MPI_Send and *MPI_Recv* for point to point communication.

MPI_ANY_SOURCE and *MPI_ANY_TAG* [4]for flexible communication.

3 Experiments and Results

3.1 Distribution of X after time T

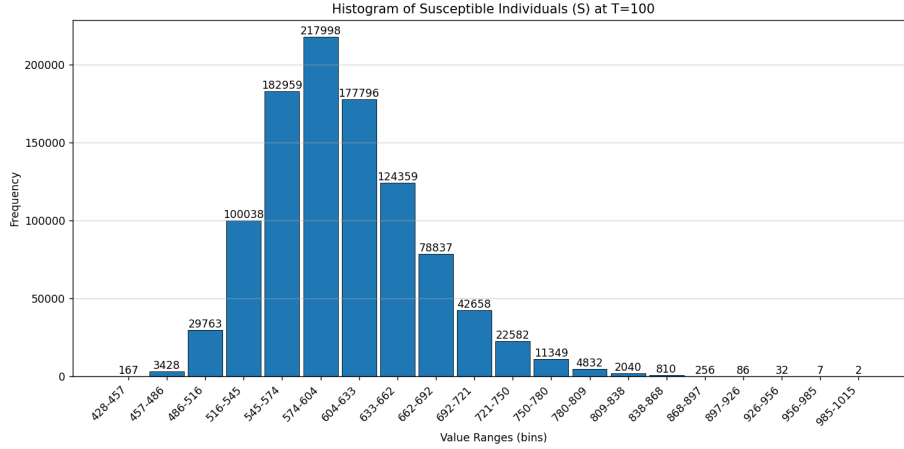


Figure 1: Distribution of X after time T

The figure above shows the histogram plot of the distribution of X at time step T given the same input $x_0 = 900, 900, 30, 330, 50, 270, 20$, and it fits our expectation that is a normal distribution and it peaks at [573,604].

4 Performance

4.1 Strong scalability

Fixed-size scalability was tested with $N_{total} = 10^6$ and varying $p = 1, 2, 4, 8, 16$.

Table 1: Strong Scalability Results

Processes Ideal Efficiency (%)	Runtime (s)	Speedup (%)	Efficiency (%)	Ideal Speedup (%)
1	1244.614	100%	100%	100%
2	1240.079	100%	50%	200%
4	446.136	279%	70%	400%
8	211.300	589%	74%	800%
16	99.826	1247%	78%	1600%




Figure 2: The Speedup in strong scalability




Figure 3: The Efficiency in strong scalability

From the result above, the runtime of 2 processes is almost equal to 1 process, but with the increase of processes (2, 4, 8, 16), we can find the improvement of speedup as we expected. For 16 processes, the runtime is reduced to 99.826 seconds, it achieves the speedup of 12.47.

4.2 weak scalability

Weak scalability was tested by increasing N_{total} proportionally with p , keeping the workload per process constant ($n = 250,000$).

Table 2: Weak Scalability Results

Processes	Runtime (s)	Speedup (%)	Ideal Speedup (%)
1	284.663	100%	100%
2	585.387	49%	100%
4	428.362	66%	100%
8	422.527	67%	100%
16	398.103	72%	100%




Figure 4: The Runtime and Speedup in strong scalability

The execution time increase significantly when moving from 1 process to 2 processes, in the work per process remaining constant, beyond 2 process, the runtime and speedup remains almost constant.

5 Discussion

In strong scalability, from 1 process to 2 processes, the runtime even not change, it indicates the overhead of communication and master-worker coordination for small number of processes. However, we can see an obvious improvement beyond 2 processes, this proof that the benefit of distributing tasks handle outweigh the

communication overhead. With the use of dynamical load balancing, the larger numbers of workers can keep them busy and reduce idle time.

In weak scalability, we met same problem, the significant increase of execution time when we moved from 1 process to 2 processes, it can be attributes to the master-worker framework, because when running on more processes, the master must begin sending tasks and receiving results, which adds the computational burden and leads to overweight the benefit of distribution. On contrast, when we run a single process, the overheads of message passing are completely disappear. But beyond 2 processes, the runtime is almost constant, and this suggests that for a larger number of processes, the dynamic load balancing becomes more efficient.

6 Conclusion

Both strong and weak scaling tests shows a significant overhead when introduced it to an MPI-parallelized master-worker architecture, and it is prominent for small number of processes. But for larger process counts(beyond 2 processes), the dynamic load balancing effectively utilizes the available cores. And there has a lot of way of improvement(e.g. a manager-worker model where the manager also participates in computation but when tried to add master(Rank 0) into computation as worker, it increases the execution time due to increase burden of rank 0).

7 Declaration of AI

In my project, I used AI to write some repeated and basic codes, debugging and make a improvement of my code. I write the original code frame and tell AI my the ideas, from the materials from AI's search, I found a way to improve performance and looked for the bibliographic on online library, it increases the efficiency a lot, I don't need to waste a lot time in searching method. AI also helped me debug my code when I met some errors on the UPPMAX computing cluster, and modifying my code based on my results, I always get some useful feedback from AI.

References

- [1] Myles R. Allen and Leonard A. Smith. Monte carlo ssa: Detecting irregular oscillations in the presence of colored noise. *Journal of climate*, 9(12):3373–3404, 1996.

- [2] Rosa Filgueira, Jesús Carretero, David E. Singh, Alejandro Calderón, and Alberto Núñez. Dynamic-compi: dynamic optimization techniques for mpi parallel applications. *The Journal of supercomputing*, 59(1):361–391, 2012.
- [3] Mats Rynge, Scott Callaghan, Ewa Deelman, Gideon Juve, Gaurang Mehta, Karan Vahi, and Philip J. Maechling. Enabling large-scale scientific workflows on petascale resources using mpi master/worker. In *Proceedings of the 1st Conference of the Extreme Science and Engineering Discovery Environment: Bridging from the eXtreme to the campus and beyond*, pages 1–8, New York, NY, USA, 2012. ACM.
- [4] Anh Vo, Sarvani Vakkalanka, Jason Williams, Ganesh Gopalakrishnan, Robert M. Kirby, Rajeev Thakur, Jan Westerholm, Jack Dongarra, and Matti Ropo. Sound and efficient dynamic verification of mpi programs with probe non-determinism. In *Recent Advances in Parallel Virtual Machine and Message Passing Interface*, Lecture Notes in Computer Science, pages 271–281. Springer Berlin Heidelberg, Berlin, Heidelberg, 2009.

8 Appendix

```

1
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <math.h>
5 #include <mpi.h>
6 #include <time.h>
7 #include <string.h>
8 #define R_ssa 15
9 #define X_LEN 7
10 #define T_FINAL 100.0
11
12
13 int P[R_ssa][X_LEN] = {
14     { 1,  0,  0,  0,  0,  0,  0},
15     {-1,  0,  0,  0,  0,  0,  0},
16     {-1,  0,  1,  0,  0,  0,  0},
17     { 0,  1,  0,  0,  0,  0,  0},
18     { 0, -1,  0,  0,  0,  0,  0},
19     { 0, -1,  0,  1,  0,  0,  0},
20     { 0,  0, -1,  0,  0,  0,  0},
21     { 0,  0, -1,  0,  1,  0,  0},
22     { 0,  0,  0, -1,  0,  0,  0},
23     { 0,  0,  0, -1,  0,  1,  0},
24     { 0,  0,  0,  0, -1,  0,  0},
25     { 0,  0,  0,  0, -1,  0,  1},
26     { 0,  0,  0,  0,  0, -1,  0},
27     { 1,  0,  0,  0,  0,  0, -1},
28     { 0,  0,  0,  0,  0,  0, -1}

```

```

29 };
30
31 void prop(int *x, double *w) {
32     const double LAMBDA_H = 20;
33     const double LAMBDA_M = 0.5;
34     const double B = 0.075;
35     const double BETA_H = 0.3;
36     const double BETA_M = 0.5;
37     const double MU_H = 0.015;
38     const double MU_M = 0.02;
39     const double DELTA_H = 0.05;
40     const double DELTA_M = 0.15;
41     const double ALFA_H = 0.6;
42     const double ALFA_M = 0.6;
43     const double R = 0.05;
44     const double OMEGA = 0.02;
45     const double NU_H = 0.5;
46     const double NU_M = 0.15;
47
48     w[0] = LAMBDA_H;
49     w[1] = MU_H * x[0];
50     w[2] = (B * BETA_H * x[0] * x[5]) / (1 + NU_H * x[5]);
51     w[3] = LAMBDA_M;
52     w[4] = MU_M * x[1];
53     w[5] = (B * BETA_M * x[1] * x[4]) / (1 + NU_M * x[4]);
54     w[6] = MU_H * x[2];
55     w[7] = ALFA_H * x[2];
56     w[8] = MU_M * x[3];
57     w[9] = ALFA_M * x[3];
58     w[10] = (MU_H + DELTA_H) * x[4];
59     w[11] = R * x[4];
60     w[12] = (MU_M + DELTA_M) * x[5];
61     w[13] = OMEGA * x[6];
62     w[14] = MU_H * x[6];
63 }
64
65 void SSA(double T, int *x0, int *x_final) {
66     double t = 0.0;
67     int x[X_LEN];
68     double w[R_ssa];
69     double a0, tau;
70     int i, r;
71     double u1, u2, sum;
72
73     for (i = 0; i < X_LEN; i++) x[i] = x0[i];
74
75     while (t < T) {
76         prop(x, w);
77         a0 = 0.0;
78         for (i = 0; i < R_ssa; i++) a0 += w[i];

```



```

79         if (a0 == 0.0) break;
80
81         u1 = (double) rand() / RAND_MAX;
82         u2 = (double) rand() / RAND_MAX;
83         tau = -log(u1) / a0;
84         t += tau;
85         if (t >= T) break;
86
87         sum = 0.0;
88         for (r = 0; r < R_ssa; r++) {
89             sum += w[r];
90             if (sum >= u2 * a0) break;
91         }
92
93         for (i = 0; i < X_LEN; i++)
94             x[i] += P[r][i];
95     }
96
97     for (i = 0; i < X_LEN; i++)
98         x_final[i] = x[i];
99 }
100
101 void build_histogram(int *values, int N, int bins) {
102     int min_val = values[0];
103     int max_val = values[0];
104     for (int i = 1; i < N; i++) {
105         if (values[i] < min_val) min_val = values[i];
106         if (values[i] > max_val) max_val = values[i];
107     }
108
109     int *hist = calloc(bins, sizeof(int));
110     double bin_width = (double)(max_val - min_val) / bins;
111
112     for (int i = 0; i < N; i++) {
113         int bin_index = (int)((values[i] - min_val) /
114                               bin_width);
115         if (bin_index == bins) bin_index = bins - 1;
116         hist[bin_index]++;
117     }
118
119     printf("Histogram bins and counts:\n");
120     for (int i = 0; i < bins; i++) {
121         double bin_start = min_val + i * bin_width;
122         double bin_end = bin_start + bin_width;
123         printf("Bin %2d: [%6.2f - %6.2f] Count: %d\n", i,
124               bin_start, bin_end, hist[i]);
125     }
126
127     free(hist);
128 }

```

```

127 void run_simulation(double *result) {
128     int x0[X_LEN] = {900, 900, 30, 330, 50, 270, 20};
129     int x_final[X_LEN];
130     SSA(T_FINAL, x0, x_final);
131     for (int i = 0; i < X_LEN; i++) {
132         result[i] = (double)x_final[i];
133     }
134 }
135
136 int main(int argc, char *argv[]) {
137     MPI_Init(&argc, &argv);
138     int rank, size;
139     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
140     MPI_Comm_size(MPI_COMM_WORLD, &size);
141     srand(time(NULL) + rank);
142     if (argc < 2) {
143         if (rank == 0) {
144             printf("Usage: %s <total_runs>\n", argv[0]);
145         }
146         MPI_Finalize();
147         return 1;
148     }
149
150     int total_runs = atoi(argv[1]);
151
152
153     double start_time = MPI_Wtime();
154
155     /*Here, I used the dynamic loading balancing, so each
156     worker gets a new tasks
157     only when it finishes the previous one, I learn this way
158     from the chagpt and the course of HPP*/
159     double *all_results = NULL;
160     if (rank == 0) {
161         all_results = malloc(total_runs * X_LEN * sizeof(
162             double));
163         if (all_results == NULL) {
164             fprintf(stderr, "Master: Failed to allocate
165             memory for results.\n");
166             MPI_Abort(MPI_COMM_WORLD, 1);
167         }
168     }
169     if (size == 1) {
170         // If only one process, run all simulations
171         sequentially on rank 0
172         if (rank == 0) {
173             printf("Running in single-process mode (rank 0
174             performs all %d simulations).\n", total_runs)
175             ;
176             for (int i = 0; i < total_runs; ++i) {

```

```

170         double result[X_LEN];
171         run_simulation(result);
172         memcpy(&all_results[i * X_LEN], result,
               X_LEN * sizeof(double));
173     }
174 }
175 } else {
176     // Master-worker with dynamic load balancing
177     if (rank == 0) {
178         int num_sent = 0;
179         int num_results = 0;
180
181         // 1. Master sends initial tasks to all workers
182         // (excluding itself for now)
183         for (int i = 1; i < size && num_sent <
              total_runs; i++) {
184             MPI_Send(&num_sent, 1, MPI_INT, i, 0,
                     MPI_COMM_WORLD);
185             num_sent++;
186         }
187
188         // 2. Master continues to receive results and
189         // send new tasks.
190         while (num_results < total_runs) {
191             double result[X_LEN];
192             MPI_Status status;
193
194             // 3. Receive results from any worker
195             MPI_Recv(result, X_LEN, MPI_DOUBLE,
                     MPI_ANY_SOURCE, MPI_ANY_TAG,
                     MPI_COMM_WORLD, &status);
196             int sender = status.MPI_SOURCE;
197             memcpy(&all_results[num_results * X_LEN],
                     result, X_LEN * sizeof(double));
198             num_results++;
199
200             // 4. Send next task if available
201             if (num_sent < total_runs) {
202                 MPI_Send(&num_sent, 1, MPI_INT, sender,
                         0, MPI_COMM_WORLD);
203                 num_sent++;
204             } else {
205                 int stop_signal = -1;
206                 MPI_Send(&stop_signal, 1, MPI_INT,
                         sender, 0, MPI_COMM_WORLD);
207             }
208         }

```

```

209         // After all results are collected, send stop
        signals to any remaining idle workers
210         for (int i = 1; i < size; ++i) {
211             int stop_signal = -1;
212             MPI_Send(&stop_signal, 1, MPI_INT, i, 0,
                     MPI_COMM_WORLD);
213         }
214
215     } else {
216         int task_index;
217         while (1) {
218             MPI_Recv(&task_index, 1, MPI_INT, 0, 0,
                     MPI_COMM_WORLD, MPI_STATUS_IGNORE);
219             if (task_index == -1) break;
220
221             double result[X_LEN];
222             run_simulation(result);
223             MPI_Send(result, X_LEN, MPI_DOUBLE, 0, 0,
                     MPI_COMM_WORLD);
224         }
225     }
226 }
227 if (rank == 0) {
228     FILE *fp = fopen("histogram_data.csv", "w");
229     if (fp != NULL) {
230         for (int i = 0; i < total_runs; i++) {
231             double *x = &all_results[i * X_LEN];
232             fprintf(fp, "%.2f,%.2f,%.2f\n", x[0], x[1],
                     x[2]);
233         }
234         fclose(fp);
235     } else {
236         printf("Error opening file for writing.\n");
237     }
238
239     int *S_values = malloc(total_runs * sizeof(int));
240     for (int i = 0; i < total_runs; i++) {
241         S_values[i] = (int) all_results[i * X_LEN];
242     }
243     build_histogram(S_values, total_runs, 20);
244     free(S_values);
245     free(all_results);
246
247     double end_time = MPI_Wtime();
248     printf("Total simulation time: %.3f seconds\n",
           end_time - start_time);
249 }
250
251 MPI_Finalize();
252 return 0;

```

